

Model Specification

We'll be working today with the wage2 dataset, which includes monthly wages of male earners along with a variety of characteristics. We'll be attempting to estimate some fairly standard wage models, but we'll also try to answer the most vexing question for many students: what variables should I put in my model?

The most important answer to that question is to use theory. Theory and previous results are our only guide—the data simply can't tell you by themselves what belongs in the model and what doesn't. However, we can use a combination of theory and applied data analysis to come up with a model that fits the data well and says something interesting about theory.

Missing Data

Before we start with all that, let's talk again about how Stata handles missing data. Let's assume that we want to estimate several nested models, first with hours, education and age, then the same model with mother's education, then the same model with father's education, then a final model with all variables. Our results look like this:

```
.
. reg lwage hours educ age
```

| Source | SS | df | MS | | | |
|----------|------------|-----|------------|-----------------|--------|--|
| Model | 21.7514568 | 3 | 7.25048559 | Number of obs = | 935 | |
| Residual | 143.904838 | 931 | .15457018 | F(3, 931) = | 46.91 | |
| Total | 165.656294 | 934 | .177362199 | Prob > F = | 0.0000 | |
| | | | | R-squared = | 0.1313 | |
| | | | | Adj R-squared = | 0.1285 | |
| | | | | Root MSE = | .39315 | |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|-------|-------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | | -.0047011 | .0017887 | -2.63 | 0.009 | -.0082115 | -.0011906 |
| educ | | .0616404 | .0058814 | 10.48 | 0.000 | .0500981 | .0731827 |
| age | | .0227339 | .0041411 | 5.49 | 0.000 | .0146069 | .0308608 |
| _cons | | 5.403279 | .1732026 | 31.20 | 0.000 | 5.063366 | 5.743191 |


```
.
. reg lwage hours educ age meduc
```

| Source | SS | df | MS | | | |
|----------|------------|-----|------------|-----------------|--------|--|
| Model | 22.8514162 | 4 | 5.71285406 | Number of obs = | 857 | |
| Residual | 126.509635 | 852 | .148485487 | F(4, 852) = | 38.47 | |
| Total | 149.361051 | 856 | .174487209 | Prob > F = | 0.0000 | |
| | | | | R-squared = | 0.1530 | |
| | | | | Adj R-squared = | 0.1490 | |
| | | | | Root MSE = | .38534 | |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|-------|-------|-------|-----------|---|------|----------------------|--|
| hours | | | | | | | |
| educ | | | | | | | |
| age | | | | | | | |
| meduc | | | | | | | |
| _cons | | | | | | | |

| | | | | | | |
|-------|-----------|----------|-------|-------|-----------|-----------|
| hours | -.0058052 | .0018374 | -3.16 | 0.002 | -.0094115 | -.0021989 |
| educ | .0525597 | .0064521 | 8.15 | 0.000 | .0398957 | .0652236 |
| age | .0243798 | .0042747 | 5.70 | 0.000 | .0159896 | .03277 |
| meduc | .0184424 | .0049725 | 3.71 | 0.000 | .0086826 | .0282022 |
| _cons | 5.33402 | .1776844 | 30.02 | 0.000 | 4.98527 | 5.682771 |

```
.
. reg lwage hours educ age feduc
```

| Source | SS | df | MS | Number of obs = | 741 |
|----------|------------|-----|------------|-----------------|----------|
| | | | | F(4, 736) = | 34.17 |
| Model | 20.2139719 | 4 | 5.05349299 | Prob > F | = 0.0000 |
| Residual | 108.836202 | 736 | .147875274 | R-squared | = 0.1566 |
| | | | | Adj R-squared | = 0.1521 |
| Total | 129.050173 | 740 | .174392126 | Root MSE | = .38455 |

| lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|-------|----------|-----------|-------|-------|----------------------|-----------|
| hours | -.007041 | .0019639 | -3.59 | 0.000 | -.0108964 | -.0031855 |
| educ | .0475258 | .0070026 | 6.79 | 0.000 | .0337785 | .0612732 |
| age | .0262759 | .0045806 | 5.74 | 0.000 | .0172834 | .0352685 |
| feduc | .0172076 | .0047569 | 3.62 | 0.000 | .0078689 | .0265462 |
| _cons | 5.421121 | .1897058 | 28.58 | 0.000 | 5.048692 | 5.79355 |

```
.
. reg lwage hours educ age feduc meduc
```

| Source | SS | df | MS | Number of obs = | 722 |
|----------|------------|-----|------------|-----------------|----------|
| Model | 20.519095 | 5 | 4.10381899 | F(5, 716) = | 27.64 |
| Residual | 106.292836 | 716 | .148453682 | Prob > F | = 0.0000 |
| | | | | R-squared | = 0.1618 |
| | | | | Adj R-squared | = 0.1560 |
| Total | 126.811931 | 721 | .1758834 | Root MSE | = .3853 |

| lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|-------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | -.0071408 | .0019821 | -3.60 | 0.000 | -.0110321 | -.0032495 |
| educ | .046006 | .0072322 | 6.36 | 0.000 | .0318071 | .0602049 |
| age | .0249456 | .0046705 | 5.34 | 0.000 | .0157762 | .034115 |
| feduc | .0114239 | .0055571 | 2.06 | 0.040 | .0005138 | .022334 |
| meduc | .0132414 | .0063094 | 2.10 | 0.036 | .0008543 | .0256285 |
| _cons | 5.404781 | .1929204 | 28.02 | 0.000 | 5.026024 | 5.783539 |

The results are extremely problematic because each set of results is on a different sample! The first set has 857 observations, the second 741, and down to 722 for the final one. Stata performs casewise deletion when running regressions, and doesn't adjust unless you tell it to. In this case none of the standard tests of model fit are relevant, because it's not the same sample.

The solution is to use the `e(sample)` command to limit the sample to the relevant analysis sample. First, run the model that restricts the data the most (has the most missing data), then limit subsequent models using the statement `if e(sample)==1`.

The natural log transformation

The variable `lwage` is the natural log of wages. This means that it has been transformed by taking the natural log of the underlying variable:

$$\log_e(y_i) = x \equiv e^x = y_i$$

Where e is Euler's constant, $e = \sum_{n=0}^{\infty} \frac{1}{n!} = \frac{1}{1} + \frac{1}{1} + \frac{1}{12} + \frac{1}{123} \dots$

The log transformation is used all the time, and particularly in econometrics. It's useful whenever you have a variable that follows some kind of exponential distribution, with widely disparate levels. Earnings, school sizes, revenues of institutions of higher education and state populations are all examples of these kinds of situations.

When the dependent variable is log transformed but the independent variable is not, this is called a log-level regression. In a log-level regression,

$$\log(y_i) = \beta_0 + \beta_1 x_i + \epsilon_i$$

Which implies that

$$y_i = e^{\beta_0 + \beta_1 x_i + \epsilon_i}$$

And . . .

$$\frac{dy}{dx} = \beta e^{\beta_0 + \beta_1 x_i + \epsilon} = \beta_1 y$$

Which means that the coefficient, β_1

$$\beta_1 = \frac{dy}{dx} \frac{1}{y}$$

This changes our interpretation to mean that for a one unit increase in x , y is predicted to increase by β_1 proportion of y or more commonly by $100 * \beta_1$ percent. It changes the scale of the dependent variable to be on the $1/y$ scale as opposed to the y scale, so everything is about a proportional (or percentage) increase in y .

Quick Exercise

Interpret the coefficients from the basic earnings regression of log wages on years of education.

Selecting Variables: Stepwise Regression OR A Cautionary Tale of Woe

When selecting variables for a model, students are sometimes tempted by the dark side of stepwise regression, which is a step on the path toward the greater evil that is data mining. I will

illustrate why this is a bad idea. The basic idea with stepwise regression is to eliminate variables from the model one at a time—if the variable is not significant, it gets dropped. However, this method is very sensitive to the overall group of variables used, essentially just pushing decisions one step back, and then using an arbitrary non-theoretical standard for variable inclusion. There is no good theoretical reason to use this procedure.

Selecting Variables: RESET test

One question that comes up frequently is whether one or more variables ought to be expressed as quadratic or higher-order polynomials in the equation. The RESET test can help with this problem. Specifying the RESET test without any options means that Stata will fit the model with the second, third and fourth powers of \hat{y} . Specifying the option `rhs` will use powers of the individual regressors.

In Stata, we would run:

```
.
. reg lwage hours age educ

      Source |      SS      df      MS                Number of obs =      935
-----+-----+-----+-----+-----+-----+-----+-----
      Model | 21.7514568      3  7.25048559             F( 3,  931) =    46.91
      Residual | 143.904838    931  .15457018             Prob > F      =    0.0000
-----+-----+-----+-----+-----+-----+-----
      Total | 165.656294    934  .177362199             R-squared     =    0.1313
                                           Adj R-squared =    0.1285
                                           Root MSE     =    .39315

-----+-----+-----+-----+-----+-----+-----
      lwage |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
-----+-----+-----+-----+-----+-----+-----
      hours | -.0047011   .0017887     -2.63   0.009   -.0082115   -.0011906
        age | .0227339   .0041411      5.49   0.000    .0146069    .0308608
        educ | .0616404   .0058814     10.48   0.000    .0500981    .0731827
       _cons | 5.403279   .1732026     31.20   0.000    5.063366    5.743191
-----+-----+-----+-----+-----+-----+-----

.
. estat ovtest

Ramsey RESET test using powers of the fitted values of lwage
Ho: model has no omitted variables
      F(3, 928) =      0.32
      Prob > F =      0.8089

.
. estat ovtest, rhs

Ramsey RESET test using powers of the independent variables
Ho: model has no omitted variables
      F(9, 922) =      2.12
      Prob > F =      0.0255
```

The result of the first test is not significant, but the result of the second test is. This indicates that we might want to include some additional powers of the right hand variables. Let's begin by introducing a quadratic function of age:

```
. gen agesq=age^2

.
. label var agesq "Age squared"
```

```

. reg lwage hours educ age agesq

```

| Source | SS | df | MS | | | |
|----------|------------|-----|------------|-----------------|--------|--|
| Model | 21.7551592 | 4 | 5.43878981 | Number of obs = | 935 | |
| Residual | 143.901135 | 930 | .154732403 | F(4, 930) = | 35.15 | |
| Total | 165.656294 | 934 | .177362199 | Prob > F = | 0.0000 | |
| | | | | R-squared = | 0.1313 | |
| | | | | Adj R-squared = | 0.1276 | |
| | | | | Root MSE = | .39336 | |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|-------|-------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | | -.0047 | .0017897 | -2.63 | 0.009 | -.0082123 | -.0011876 |
| educ | | .0615585 | .0059083 | 10.42 | 0.000 | .0499634 | .0731536 |
| age | | .0388675 | .1043805 | 0.37 | 0.710 | -.1659811 | .2437162 |
| agesq | | -.0002425 | .0015678 | -0.15 | 0.877 | -.0033193 | .0028343 |
| _cons | | 5.138356 | 1.721371 | 2.99 | 0.003 | 1.760134 | 8.516578 |


```

. test age agesq

( 1) age = 0
( 2) agesq = 0

F( 2, 930) = 15.07
Prob > F = 0.0000

```

The two terms for age are jointly significant, but it looks like we could safely exclude age squared from the model without any loss of model fit.

Now let's try education squared:

```

. gen educsq=educ^2

. la var educsq "Education squared"

. reg lwage hours age educ educsq

```

| Source | SS | df | MS | | | |
|----------|------------|-----|------------|-----------------|--------|--|
| Model | 22.3576551 | 4 | 5.58941378 | Number of obs = | 935 | |
| Residual | 143.298639 | 930 | .154084558 | F(4, 930) = | 36.27 | |
| Total | 165.656294 | 934 | .177362199 | Prob > F = | 0.0000 | |
| | | | | R-squared = | 0.1350 | |
| | | | | Adj R-squared = | 0.1312 | |
| | | | | Root MSE = | .39254 | |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|--------|-------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | | -.004608 | .0017866 | -2.58 | 0.010 | -.0081141 | -.0011018 |
| age | | .0243563 | .0042147 | 5.78 | 0.000 | .0160848 | .0326278 |
| educ | | .2161619 | .0781252 | 2.77 | 0.006 | .0628397 | .369484 |
| educsq | | -.0054912 | .0027685 | -1.98 | 0.048 | -.0109245 | -.000058 |
| _cons | | 4.286926 | .5887928 | 7.28 | 0.000 | 3.13141 | 5.442443 |

This does result in a statistically significant increase in model fit. The way I would prefer approaching this problem is to fully specify the model, then restrict it appropriately, like so:

```

. /*Preferred method */

```

```

. reg lwage hours age agesq educ educsq

```

| Source | SS | df | MS | | | |
|----------|------------|-----|------------|-----------------|--------|--|
| Model | 22.3674226 | 5 | 4.47348452 | Number of obs = | 935 | |
| Residual | 143.288872 | 929 | .154239905 | F(5, 929) = | 29.00 | |
| Total | 165.656294 | 934 | .177362199 | Prob > F = | 0.0000 | |
| | | | | R-squared = | 0.1350 | |
| | | | | Adj R-squared = | 0.1304 | |
| | | | | Root MSE = | .39273 | |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|--------|-------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | | -.0046056 | .0017875 | -2.58 | 0.010 | -.0081135 | -.0010976 |
| age | | .050602 | .1043806 | 0.48 | 0.628 | -.154247 | .255451 |
| agesq | | -.0003944 | .0015672 | -0.25 | 0.801 | -.0034699 | .0026812 |
| educ | | .2169837 | .0782328 | 2.77 | 0.006 | .0634502 | .3705171 |
| educsq | | -.0055252 | .0027732 | -1.99 | 0.047 | -.0109676 | -.0000828 |
| _cons | | 3.849225 | 1.836393 | 2.10 | 0.036 | .2452658 | 7.453184 |


```

. test age agesq

```

(1) age = 0
 (2) agesq = 0

F(2, 929) = 16.71
 Prob > F = 0.0000


```

. test educ educsq

```

(1) educ = 0
 (2) educsq = 0

F(2, 929) = 56.44
 Prob > F = 0.0000

Selecting Variables: Non-Nested Models

In many situations, models are based on competing hypotheses, and so they don't nest within one another. Let's say we have one model that posits education as the key to wages, another that posits iq as the key to wages. To test whether one is better than the other, we use the Davidson-Mackinnon test:

```

. reg lwage hours iq

```

| Source | SS | df | MS | | | |
|----------|------------|-----|------------|-----------------|--------|--|
| Model | 17.2420918 | 2 | 8.62104588 | Number of obs = | 935 | |
| Residual | 148.414203 | 932 | .159242707 | F(2, 932) = | 54.14 | |
| Total | 165.656294 | 934 | .177362199 | Prob > F = | 0.0000 | |
| | | | | R-squared = | 0.1041 | |
| | | | | Adj R-squared = | 0.1022 | |
| | | | | Root MSE = | .39905 | |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|-------|-------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | | -.0041302 | .0018124 | -2.28 | 0.023 | -.007687 | -.0005734 |
| iq | | .0089535 | .0008698 | 10.29 | 0.000 | .0072465 | .0106606 |
| _cons | | 6.053607 | .1150416 | 52.62 | 0.000 | 5.827837 | 6.279378 |

```
.
. reg lwage hours educ
```

| Source | SS | df | MS | | Number of obs = | 935 |
|----------|------------|-----|------------|--|-----------------|--------|
| Model | 17.0930188 | 2 | 8.54650938 | | F(2, 932) = | 53.62 |
| Residual | 148.563276 | 932 | .159402656 | | Prob > F = | 0.0000 |
| | | | | | R-squared = | 0.1032 |
| | | | | | Adj R-squared = | 0.1013 |
| Total | 165.656294 | 934 | .177362199 | | Root MSE = | .39925 |

| | lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] |
|-------|-------|-----------|-----------|-------|-------|----------------------|
| hours | | -.0044454 | .0018159 | -2.45 | 0.015 | -.0080091 -.0008817 |
| educ | | .0611697 | .005972 | 10.24 | 0.000 | .0494497 .0728898 |
| _cons | | 6.150426 | .108791 | 56.53 | 0.000 | 5.936923 6.36393 |


```
.
. nnest lwage hours iq (hours educ)

M1 : Y = a + Xb with X = [hours iq]
M2 : Y = a + Zg with Z = [hours educ]

J test for non-nested models

H0 : M1 t(931)      5.89476
H1 : M2 p-val      0.00000

H0 : M2 t(931)      5.97647
H1 : M1 p-val      0.00000

Cox-Pesaran test for non-nested models

H0 : M1 N(0,1)      -8.87744
H1 : M2 p-val      0.00000

H0 : M2 N(0,1)      -9.05449
H1 : M1 p-val      0.00000

.
```

The results of this test indicate that it would be better to include both of these models, in a sort of “super” model.

Interactions

Interactions can be difficult to understand, but they are key to getting a handle on sometimes important moderating variables in an analysis.

Interactions with two binary variables

Let’s say we’re interested in whether marriage affects wages differently for black and white men. The specification of an interaction between the two binary variables of white and married would look like this:

```
. reg lwage hours age educ i.black#i.married iq meduc south urban
```

| Source | SS | df | MS | Number of obs = 857 | | |
|----------|------------|-----|------------|------------------------|--|--|
| Model | 38.9381574 | 10 | 3.89381574 | F(10, 846) = 29.83 | | |
| Residual | 110.422894 | 846 | .130523515 | Prob > F = 0.0000 | | |
| Total | 149.361051 | 856 | .174487209 | R-squared = 0.2607 | | |
| | | | | Adj R-squared = 0.2520 | | |
| | | | | Root MSE = .36128 | | |

| lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|---------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | -.0063301 | .0017273 | -3.66 | 0.000 | -.0097204 | -.0029398 |
| age | .0220385 | .0040456 | 5.45 | 0.000 | .0140979 | .0299792 |
| educ | .0363765 | .0068217 | 5.33 | 0.000 | .0229871 | .0497659 |
| black# | | | | | | |
| married | | | | | | |
| 0 1 | .1817023 | .0438419 | 4.14 | 0.000 | .0956506 | .267754 |
| 1 0 | -.2750052 | .1036788 | -2.65 | 0.008 | -.478503 | -.0715073 |
| 1 1 | .0597192 | .0605138 | 0.99 | 0.324 | -.0590557 | .178494 |
| iq | .0039649 | .0010438 | 3.80 | 0.000 | .0019162 | .0060137 |
| meduc | .0104632 | .0047967 | 2.18 | 0.029 | .0010483 | .0198781 |
| south | -.0729291 | .0276262 | -2.64 | 0.008 | -.1271531 | -.0187051 |
| urban | .179466 | .0280946 | 6.39 | 0.000 | .1243228 | .2346092 |
| _cons | 5.084363 | .1866763 | 27.24 | 0.000 | 4.717961 | 5.450766 |

Quick Exercise

Run a regression with an interaction between urban and south. Interpret the results.

Interactions with a binary and continous variable

Let's say we're interested in whether education affects wages differently for black and white men. If possible, we should start by plotting the data to see if these patterns are evident.

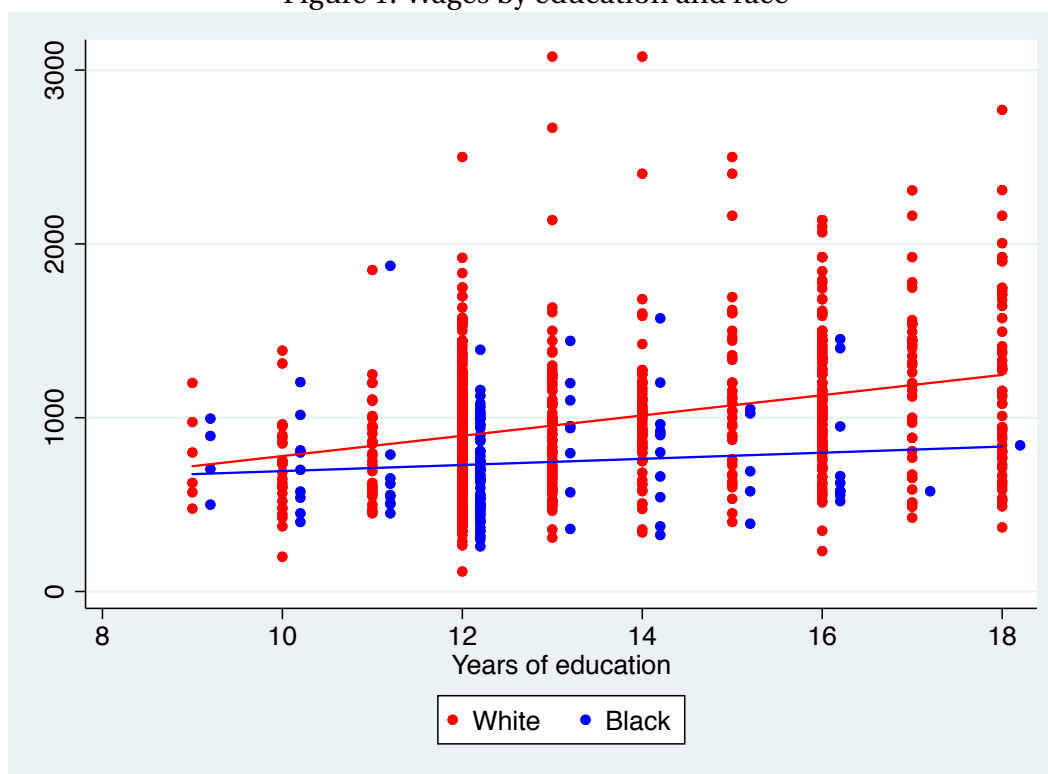
The specification of an interaction between a binary variable and a continous variable would look like this:

```
. reg lwage hours age educ black i.black#c.educ married iq meduc south urban
```

| Source | SS | df | MS | Number of obs = 857 | | |
|----------|------------|-----|------------|------------------------|--|--|
| Model | 38.7974052 | 10 | 3.87974052 | F(10, 846) = 29.69 | | |
| Residual | 110.563646 | 846 | .130689889 | Prob > F = 0.0000 | | |
| Total | 149.361051 | 856 | .174487209 | R-squared = 0.2598 | | |
| | | | | Adj R-squared = 0.2510 | | |
| | | | | Root MSE = .36151 | | |

| lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|--------------|-----------|-----------|-------|-------|----------------------|----------|
| hours | -.0063045 | .00173 | -3.64 | 0.000 | -.0097001 | -.002909 |
| age | .0215072 | .0040455 | 5.32 | 0.000 | .0135669 | .0294476 |
| educ | .0378404 | .0069883 | 5.41 | 0.000 | .0241239 | .0515569 |
| black | .1034702 | .2766783 | 0.37 | 0.709 | -.4395862 | .6465265 |
| black#c.educ | | | | | | |
| 1 | -.0196117 | .0215854 | -0.91 | 0.364 | -.0619789 | .0227554 |
| married | .2051596 | .0402851 | 5.09 | 0.000 | .1260892 | .28423 |
| iq | .0040191 | .0010442 | 3.85 | 0.000 | .0019695 | .0060687 |

Figure 1: Wages by education and race



| | | | | | | | |
|-------|--|----------|----------|-------|-------|-----------|-----------|
| meduc | | .0102747 | .0047964 | 2.14 | 0.032 | .0008604 | .0196889 |
| south | | -.069586 | .0276172 | -2.52 | 0.012 | -.1237922 | -.0153798 |
| urban | | .179133 | .0281112 | 6.37 | 0.000 | .1239572 | .2343088 |
| _cons | | 5.05532 | .187864 | 26.91 | 0.000 | 4.686586 | 5.424055 |

Quick Exercise

Run a regression with an interaction between urban and education. Interpret the results.

Interactions with two continuous variables

Finally let's say we think that education will affect your wages differently depending on your age. The specification of an interaction between two continuous variables would look like this:

```
. reg lwage hours age educ c.age#c.educ black married iq meduc south urban
```

| Source | | SS | df | MS | | Number of obs = | 857 |
|----------|--|------------|-----|------------|--|-----------------|--------|
| Model | | 39.1769054 | 10 | 3.91769054 | | F(10, 846) = | 30.08 |
| Residual | | 110.184146 | 846 | .130241307 | | Prob > F = | 0.0000 |
| Total | | 149.361051 | 856 | .174487209 | | R-squared = | 0.2623 |
| | | | | | | Adj R-squared = | 0.2536 |
| | | | | | | Root MSE = | .36089 |

| lwage | Coef. | Std. Err. | t | P> t | [95% Conf. Interval] | |
|--------------|-----------|-----------|-------|-------|----------------------|-----------|
| hours | -.0066151 | .0017294 | -3.83 | 0.000 | -.0100095 | -.0032208 |
| age | -.0280105 | .0260076 | -1.08 | 0.282 | -.0790575 | .0230364 |
| educ | -.0876756 | .0645406 | -1.36 | 0.175 | -.2143541 | .0390029 |
| c.age#c.educ | .0037019 | .0019136 | 1.93 | 0.053 | -.0000542 | .0074579 |
| black | -.1466181 | .0430343 | -3.41 | 0.001 | -.2310847 | -.0621515 |
| married | .2063299 | .0402134 | 5.13 | 0.000 | .1274003 | .2852596 |
| iq | .0040003 | .0010423 | 3.84 | 0.000 | .0019545 | .0060461 |
| meduc | .0100804 | .0047844 | 2.11 | 0.035 | .0006897 | .019471 |
| south | -.06698 | .0276104 | -2.43 | 0.015 | -.121173 | -.012787 |
| urban | .182844 | .02813 | 6.50 | 0.000 | .1276312 | .2380569 |
| _cons | 6.747921 | .8848456 | 7.63 | 0.000 | 5.011171 | 8.484672 |

The do file has extensive examples of how to use margins to create nice plots of various types of interactions. We'll go over these in class.

Not-so-quick Exercise

In pairs, I would like you to estimate the best possible model using the wage2 dataset. Think about model specification and functional form, with an eye toward possible non-linearities and other issues. Generate a do file that walks through your process of identifying the best model. Generate a fancy graph that shows the predictions made by your model.