

Working With Panel Data

Introduction

Panel data refers to data with multiple observations per unit. In education settings panel data is almost more common than not, with many studies involving cases that have been observed over time.

For all of the models below, I'll use the following notation:

y_{it} is the dependent variable for unit i ($i = 1 \dots n$) in time period t ($t = 1 \dots t$).

x_{it} is an independent variable for unit i at time t .

β is a coefficient on the variable x

ϵ_{it} is an error term

The terminology around panel data can be confusing, because economists and education experts discuss the same things using different names. Here's some terminology:

Panel data: when used by economists, this typically refers to a dataset where there are many more units than observations over time.

Cross-sectional time-series data: this refers to data where there are much longer time series, and fewer data points.

Hierarchical or "grouped" data: this refers to data where the observations are naturally grouped, e.g. students in classrooms, classrooms in schools. This type of data can also include multiple observations over time.

Fixed effects: when used by economists, this refers to models where the group mean is controlled for, either by subtracting it from the dependent variable or by individually controlling for each group effect via dummy variables. Also known as LSDV: least squares dummy variables. When HLM people say fixed effects, they're referring to coefficients that don't vary across groups. This is also known as a "no pooling" model.

Random effects: when used by economists, this refers to a model that allows one or more coefficients to have its own distribution with an error term. A random effects model is functionally equivalent to a Hierarchical Linear Model, although HLM imposes additional assumptions.

Describing Panel Data

The data we'll be using come from my dissertation, which prediction appropriations, tuition and financial aid at the state level using various characteristics of the political and higher education system. The data are a balanced panel of 49 states (excluding Alaska) over 16 years, 1984-1999.

To get Stata to recognize this as panel data, we need to use the `xtset` command.

```
. /* Set up data as panel data */  
. xtset state year, yearly  
      panel variable:  state (strongly balanced)  
      time variable:   year, 1984 to 1999  
      delta:          1 year
```

I tend to use two basic methods for describing panel data. First, I like to do line graphs for all of the continuous variables, which give you a very clear sense of variation across units and any time trends. It's also a good way to find data problems:

```
. xtline approps_i
```




Figure 1: Trend in Appropriations Per Student, by State

```
. xtline pub4tuit_i
```




Figure 2: Trend in Public Four-Year Tuition, by State

The other graph I like to use is a boxplot for the variable by state. This gives an excellent sense of variability both across and within units.

```
.
. #delimit ;
delimiter now ;
. graph hbox pub4tuit_i,
> over(state, sort(1) descending label(labsize(tiny) ))
>
> ;

. #delimit ;
delimiter now ;
. graph hbox approps,
> over(state, sort(1) descending label(labsize(tiny) ))
>
> ;
```

When reporting descriptives for a panel dataset, don't just give the grand mean. Provide averages and standard deviations for a subset of time periods, along with graphics similar to the above.




Figure 3: Variation in Appropriations Per Student, by State

Ordinary Least Squares

The OLS estimate for panel data is:

$$y_{it} = \alpha + \beta x_{it} + \epsilon_{it}$$

In Stata:

```
.
. local y approps_i

.
. local controls perc1824 incpcp_i percpriv taxcpc_i legcomp_i i.board

.
. reg `y' legideo `controls'
```

Source	SS	df	MS
Model	830964034	10	83096403.4
Residual	578657976	773	748587.29
Total	1.4096e+09	783	1800283.54

Number of obs = 784
F(10, 773) = 111.00
Prob > F = 0.0000
R-squared = 0.5895
Adj R-squared = 0.5842
Root MSE = 865.21

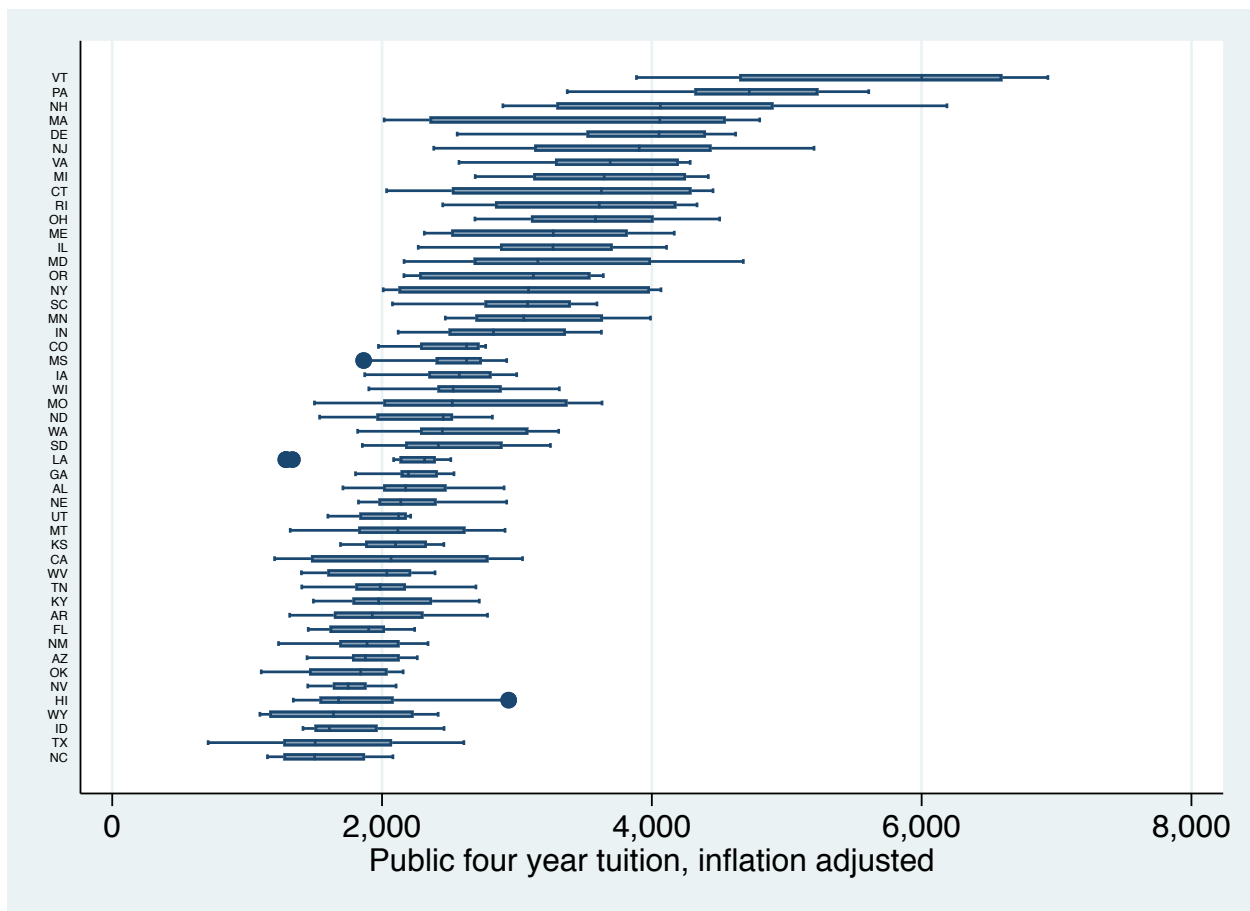


Figure 4: Variation in Public Four-Year Tuition, by State

approps_i	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
legideo	1.954075	1.373991	1.42	0.155	-.7431209	4.651272
perc1824	267.1039	29.04555	9.20	0.000	210.0865	324.1214
incpcp_i	-12.65837	12.94071	-0.98	0.328	-38.06147	12.74473
percpriv	-57.78148	2.810894	-20.56	0.000	-63.29937	-52.26359
taxcpc_i	1.939732	.1145756	16.93	0.000	1.714816	2.164649
legcomp_i	-.0008065	.0020159	-0.40	0.689	-.0047638	.0031508
board						
2	110.3047	100.8765	1.09	0.275	-87.71972	308.3291
3	-28.19471	94.82565	-0.30	0.766	-214.341	157.9516
4	-29.0085	87.02584	-0.33	0.739	-199.8435	141.8265
5	-1538.795	143.7325	-10.71	0.000	-1820.947	-1256.643
_cons	944.5651	466.3937	2.03	0.043	29.01674	1860.113

The problem with the OLS model is both that it may be inconsistent and that it may induce huge problems with heteroscedasticity. If you're not sure if you there's a problem, try graphing the residuals like so:

```
. predict e, resid
```

```
. graph box e, over(state, sort(1) descending label(labsize(tiny))) /*Horrible*/
```




Figure 5: Residuals by State

In our case, there are massive problems with the error terms by state. It's not so bad by year. Even so, we will have a correlation with the independent variables and the error term because we're leaving out a variable that is known to impact the dependent variable: the group that each unit is in.

Fixed Effects Models

The fixed effects model with group specific intercepts is:

$$y_{it} = \alpha_i + \beta x_{it} + \epsilon_{it}$$

A basic fixed effects model looking at the effect of a more liberal government on appropriations would be specified as:




Figure 6: Residuals by Year

```
. xi: xtreg `y' legideo `controls', fe
i.board      _Iboard_1-5      (naturally coded; _Iboard_1 omitted)
note: _Iboard_5 omitted because of collinearity

Fixed-effects (within) regression              Number of obs   =       784
Group variable: state                        Number of groups  =        49

R-sq:  within = 0.2281                      Obs per group: min =        16
       between = 0.0860                      avg           =       16.0
       overall  = 0.1015                      max           =        16

corr(u_i, Xb)  = -0.2562                     F(9,726)         =       23.83
                                              Prob > F          =       0.0000
```

approps_i	Coef.	Std. Err.	t	P> t	[95% Conf. Interval]	
legideo	3.508206	1.188178	2.95	0.003	1.175531	5.840881
perc1824	269.2799	24.15645	11.15	0.000	221.8551	316.7048
incpcp_i	12.85631	20.06298	0.64	0.522	-26.53207	52.24468
percpriv	-3.949699	10.39503	-0.38	0.704	-24.35761	16.45821
taxpc_i	1.436178	.1566408	9.17	0.000	1.128655	1.743701
legcomp_i	.0022197	.001997	1.11	0.267	-.0017009	.0061402
_Iboard_2	-41.50431	108.1897	-0.38	0.701	-253.9063	170.8977
_Iboard_3	-597.6449	204.1981	-2.93	0.004	-998.5342	-196.7557
_Iboard_4	-942.1278	183.3558	-5.14	0.000	-1302.099	-582.1569
_Iboard_5	(omitted)					
_cons	-38.59943	659.1182	-0.06	0.953	-1332.605	1255.406

```

-----+-----
      sigma_u | 1232.7514
      sigma_e | 492.51715
          rho | .86235025   (fraction of variance due to u_i)
-----+-----
F test that all u_i=0:      F(48, 726) =    34.57          Prob > F
= 0.0000

. xi: reg `y' legideo `controls' i.state
i.board      _Iboard_1-5      (naturally coded; _Iboard_1 omitted)
i.state      _Istate_2-50     (naturally coded; _Istate_2 omitted)
note: _Istate_22 omitted because of collinearity

      Source |      SS      df      MS              Number of obs =    784
-----+-----
      Model | 1.2335e+09    57   21640594.9      F( 57, 726) =    89.21
      Residual | 176108102   726   242573.144      Prob > F      =    0.0000
-----+-----
      Total | 1.4096e+09   783   1800283.54      R-squared     =    0.8751
                                          Adj R-squared =    0.8653
                                          Root MSE     =    492.52

-----+-----
      approps_i |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
-----+-----
      legideo | 3.508206   1.188178     2.95   0.003     1.175531     5.840881
      perc1824 | 269.2799   24.15645    11.15   0.000     221.8551     316.7048
      incpcp_i | 12.85631   20.06298     0.64   0.522    -26.53207     52.24468
      percpriv | -3.949699   10.39503    -0.38   0.704    -24.35761     16.45821
      taxcpc_i | 1.436178   .1566408     9.17   0.000     1.128655     1.743701
      legcomp_i | .0022197   .001997     1.11   0.267    -.0017009     .0061402
      _Iboard_2 | -41.50431   108.1897    -0.38   0.701    -253.9063     170.8977
      _Iboard_3 | -597.6449   204.1981    -2.93   0.004    -998.5342    -196.7557
      _Iboard_4 | -942.1278   183.3558    -5.14   0.000    -1302.099    -582.1569
      _Iboard_5 | -1876.622   218.9806    -8.57   0.000    -2306.533    -1446.711
      _Istate_3 | 276.3964   184.0882     1.50   0.134    -85.01223     637.8051
      _Istate_4 | -933.8803   259.3431    -3.60   0.000    -1443.032
      -424.7282

.....

```

This includes both the standard `xtreg` command and a `reg` command, with `xi` specified to control for state level effects. The coefficients are the same. The interpretation of a fixed effects model always refers only to within-unit changes in both the independent and dependent variables.

Without correcting for time in the above model, we could introduce serially correlated error terms.

Fixed Effects for Time

In addition to specifying fixed effects for groups, the simplest approach to handling time is to specify fixed effects for time, with T-1 variables for time included in the model, with a new set of coefficients γ_t .

$$y_{it} = \alpha_i + \beta x_{it} + \gamma_t + \epsilon_{it}$$

To estimate the above in stata, we would need to use the `i` function, which transforms variables into a categorical variable. The following syntax gives fixed effects for time, with time as a categorical variable:


```

. /* Fixed Effects for Units and Time (state and year) */
.
. xtreg `y' legideo `controls' i.year , fe
note: 5.board omitted because of collinearity

Fixed-effects (within) regression              Number of obs   =       784
Group variable: state                        Number of groups  =        49

R-sq:  within = 0.3942                      Obs per group: min =        16
        between = 0.0321                      avg           =       16.0
        overall = 0.0576                      max           =        16

corr(u_i, Xb) = -0.4822                      F(24,711)        =       19.27
                                                Prob > F          =       0.0000

-----+-----
      approps_i |          Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
-----+-----
      legideo |    1.145978    1.140491     1.00   0.315    -1.093154    3.385111
    perc1824 |   44.91926    30.75181     1.46   0.145   -15.45595   105.2945
    incpcp_i |  139.2413    28.45662     4.89   0.000    83.37225   195.1104
    percpriv |  -3.777036    10.16795    -0.37   0.710   -23.73984    16.18577
    taxcpc_i |   1.501035    .1425019    10.53   0.000     1.22126    1.78081
    legcomp_i |   .0016732    .0018213     0.92   0.359    -.0019025    .005249
      board |
    cbweak |  -24.88159    97.00964    -0.26   0.798   -215.3412   165.578
    gball |  -454.4452    183.9582    -2.47   0.014   -815.6115  -93.27897
    gbfour |  -711.7997    165.6897    -4.30   0.000   -1037.099  -386.5002
      plan |           0 (omitted)
      year |
    1985 |   220.3299    90.99408     2.42   0.016    41.68063   398.9791
    1986 |   113.4229     97.881     1.16   0.247   -78.74752   305.5932
    1987 |  -19.38342    105.0826    -0.18   0.854   -225.6928   186.926
    1988 |  -67.57019    113.1309    -0.60   0.551   -289.6807   154.5403
    1989 |  -274.3213    122.1241    -2.25   0.025   -514.0883  -34.55431
    1990 |  -399.1526    125.7373    -3.17   0.002   -646.0135  -152.2917
    1991 |  -657.3481    125.0003    -5.26   0.000   -902.7619  -411.9343
    1992 |  -678.0808    134.1735    -5.05   0.000   -941.5044  -414.6571
    1993 |   -936.106    136.6561    -6.85   0.000   -1204.404  -667.8083
    1994 |  -968.5213    145.4102    -6.66   0.000   -1254.006  -683.0365
    1995 |  -1031.559    153.3461    -6.73   0.000   -1332.624  -730.4935
    1996 |  -1044.886    162.4511    -6.43   0.000   -1363.827  -725.9445
    1997 |  -1058.236    171.5086    -6.17   0.000   -1394.96  -721.5123
    1998 |  -1197.384    189.6214    -6.31   0.000   -1569.669  -825.0989
    1999 |  -1194.228    195.1562    -6.12   0.000   -1577.379  -811.0763
      _cons |  -163.0829    776.2306    -0.21   0.834   -1687.061  1360.895
-----+-----
      sigma_u |   1421.003
      sigma_e |   440.90254
        rho |   .91218326   (fraction of variance due to u_i)
-----+-----
F test that all u_i=0: F(48, 711) = 52.77                      Prob > F = 0.0000

```

The interpretation of this would be as usual for a categorical variable: each coefficient for time represents a contrast to a base time period (stata will choose the first one). Having done this however, concerns about serial correlation should be adequately addressed.

This can be observed by looking at a boxplot of errors by time period:

```
predict res,e graph box res, over(year)
```

Fixed effects for time are not symmetric with fixed effects for groups in this model. To adjust

for this, we can regress

$$y_{*it} = y_{it} - \bar{y}_i - \bar{y}_t + \bar{y}$$

on the independent variable x , specified as:

$$x_{*it} = x_{it} - \bar{x}_i - \bar{x}_t + \bar{x}$$

Serially Correlated Errors

Fixed effects for time is an appropriate approach in many cases, however it is very inefficient: if time itself is not of interest, you will have $T - 1$ nuisance parameters along with $n - 1$ group estimates in the case of a fixed effects approach.

When estimating models for panel data, corrections for autocorrelation are much the same as in a single sample. First, assume that there is no cross-sectional autocorrelation:

$$\text{Corr}[\epsilon_{it}, \epsilon_{js}] = 0, \text{ if } i \neq j$$

In the presence of within-unit autocorrelation, the observed error ϵ_{it} consists of two parts: the error term in the previous year multiplied by a coefficient ρ and the overall error term μ_{it} .

$$\epsilon_{it} = \rho_i \epsilon_{it-1} + \mu_{it}$$

The variance of these group-specific error terms is therefore:

$$\text{Var}[\epsilon_{it}] = \sigma_i^2 = \frac{\sigma_{\mu}^2}{1 - \rho_i^2}$$

To account for this, we need to calculate a correlation coefficient ρ for each group. A group specific estimate r_i for ρ is:

$$r_i = \frac{\sum_{t=2}^T e_{it} e_{i,t-1}}{\sum_{t=1}^T e_{it}^2}$$

Most programs, including STATA, calculate a single value, which is the average of all group specific correlation coefficients. This value is then used to transform the data to eliminate the autocorrelation. For instance for y_{it} , the transformation is:

$$y_{i1}, y_{i2}, \dots, y_{iT} = \sqrt{1 - r^2} y_{i1}, y_{i2} - r_i y_{i1}, y_{i3} - r_i y_{i2}, y_{iT} - r_i y_{i,T-1}$$

To estimate a fixed effects model in STATA, use the `xtregar` command. In our running example, this can be estimated via:

```

. xi: xtregar `y' legideo `controls', fe rhtype (tsc) twostep lbi
i.board      _Iboard_1-5      (naturally coded; _Iboard_1 omitted)
note: _Iboard_5 dropped because of collinearity

FE (within) regression with AR(1) disturbances   Number of obs   =       735
Group variable: state                           Number of groups  =        49

R-sq:  within = 0.1475                          Obs per group: min =        15
       between = 0.0584                          avg =       15.0
       overall = 0.0235                          max =        15

                                           F(9,677)      =       13.02
corr(u_i, Xb) = -0.6379                      Prob > F       =       0.0000

-----+-----
      approps_i |      Coef.   Std. Err.      t    P>|t|     [95% Conf. Interval]
-----+-----
      legideo |  2.203477   1.328746     1.66   0.098   - .4054819   4.812437
    perc1824 | 308.6154   33.12758     9.32   0.000   243.5702   373.6605
    incpcp_i | 16.72646   23.96317     0.70   0.485  -30.32461   63.77753
    percpriv | 39.45439   12.99704     3.04   0.002   13.93505   64.97374
    taxcpc_i |  .9035681   .1776652     5.09   0.000   .5547272   1.252409
    legcomp_i |  .0015629   .0019213     0.81   0.416   -.0022095   .0053352
    _Iboard_2 | -88.92507   136.4487    -0.65   0.515  -356.8385   178.9884
    _Iboard_3 | -444.8078   257.5141    -1.73   0.085  -950.4301   60.8145
    _Iboard_4 | -797.8778   234.5891    -3.40   0.001  -1258.488  -337.268
    _Iboard_5 | (omitted)
      _cons | -630.4428   486.6708    -1.30   0.196  -1586.008   325.1228
-----+-----
      rho_ar |  .38558457
      sigma_u | 1626.9651
      sigma_e | 422.42905
      rho_fov |  .93684349   (fraction of variance because of u_i)
-----+-----
F test that all u_i=0:      F(48,677) =      22.21      Prob > F = 0.0000
modified Bhargava et al. Durbin-Watson = 1.0483739
Baltagi-Wu LBI = 1.2288309

```

However, the transformation of the data in the above is done via the Cochrane-Orcutt, not Prais-Winsten transformation. Cochrane-Orcutt throws out the first unit in each time series, which can be a lot of data in a panel data setting. Another option is to use xtpcse, with correlation set to AR(1). This also incorporates some other assumptions, which can be turned off by specifying the “independent” option. In particular, this allows for unit-specific autocorrelation, which is generally a better assumption.

```

. // Unit-specific ar(1) process
. xtpcse `y' legideo `controls' i.state, correlation (psar1) independent
note: 46.state omitted because of collinearity
(note: estimates of rho outside [-1,1] bounded to be in the range [-1,1])

Prais-Winsten regression, independent panels corrected standard errors

Group variable:   state                      Number of obs   =       784
Time variable:   year                        Number of groups  =        49
Panels:          independent (balanced)      Obs per group:
Autocorrelation: panel-specific AR(1)         min =        16
                                              avg =        16
                                              max =        16

Estimated covariances      =        1      R-squared      =       0.9520
Estimated autocorrelations =        49      Wald chi2(56)   =       2177.81
Estimated coefficients      =        57      Prob > chi2    =       0.0000

-----+-----
      approps_i |      Indep-corrected
      Coef.   Std. Err.      z    P>|z|     [95% Conf. Interval]
-----+-----

```

legideo		1.941346	1.113935	1.74	0.081	-.2419273	4.124619
perc1824		263.31	28.58896	9.21	0.000	207.2767	319.3434
incpcp_i		124.4886	21.42011	5.81	0.000	82.50597	166.4712
percpriv		1.636201	10.40599	0.16	0.875	-18.75916	22.03156
taxcpc_i		.5548492	.1564309	3.55	0.000	.2482503	.8614482
legcomp_i		.0029899	.0017338	1.72	0.085	-.0004084	.0063881
board							
cbweak		-103.2285	116.8174	-0.88	0.377	-332.1863	125.7294
gball		-1015.286	183.3439	-5.54	0.000	-1374.634	-655.9389
gbfour		-809.2099	194.6451	-4.16	0.000	-1190.707	-427.7125
plan		-4170.79	401.9689	-10.38	0.000	-4958.635	-3382.945

Random Effects

In the random effects model, the group effect is assumed to have a distribution and an error term. You'll get a LOT more on this in Regression II, so today I'll just introduce it to you and show you how to run the Hausman test. In practice, a random effect model is rarely appropriate unless the groups are defined as part of the sampling procedure.

First Differenced Model

First differencing is just what it sounds like: subtracting the previous time period's observation from the current one:

$$\Delta y_{it} = y_{it} - y_{it-1}$$

A first differenced model can be used with panel data, although the interpretation of coefficients goes from change in x to change-in-change in x.

$$\Delta y_{it} = \beta_0 + \beta_1 \Delta x_{it} + \epsilon$$